Full−dataset expression backend comparison

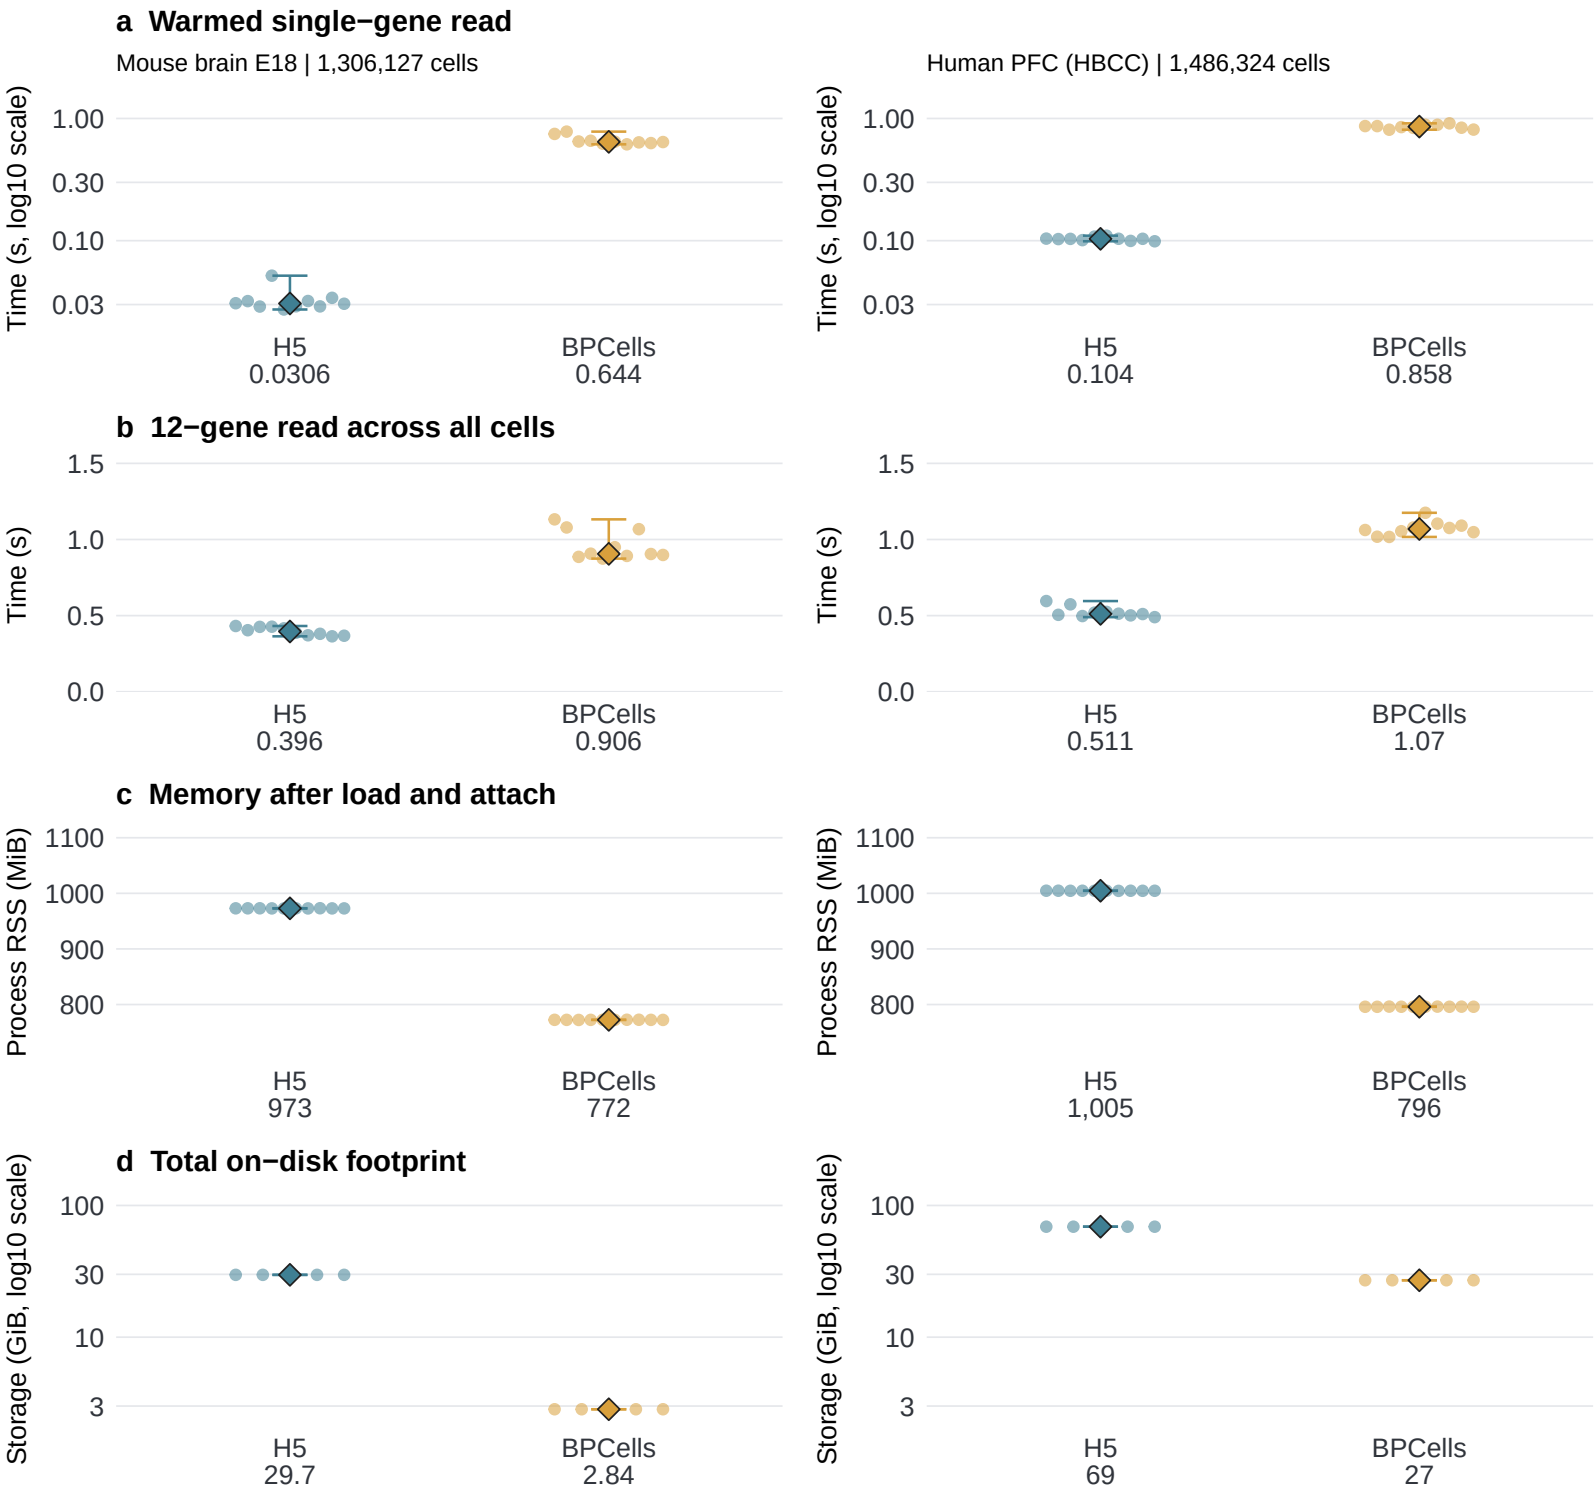

Small dots: all recorded processes. Diamonds: median. Whiskers: observed min−max (not confidence intervals).
Per dataset/backend: 5 independent builds; 10 read processes (2 per build). Read processes share built artifacts.
Values below backend names are medians. A and D: log10 axes; C: expanded linear axis (not from zero).
Single−gene timing uses each process's warmed median. Embedded exceeds sparse−index capacity.