Expression backend scaling through the full datasets

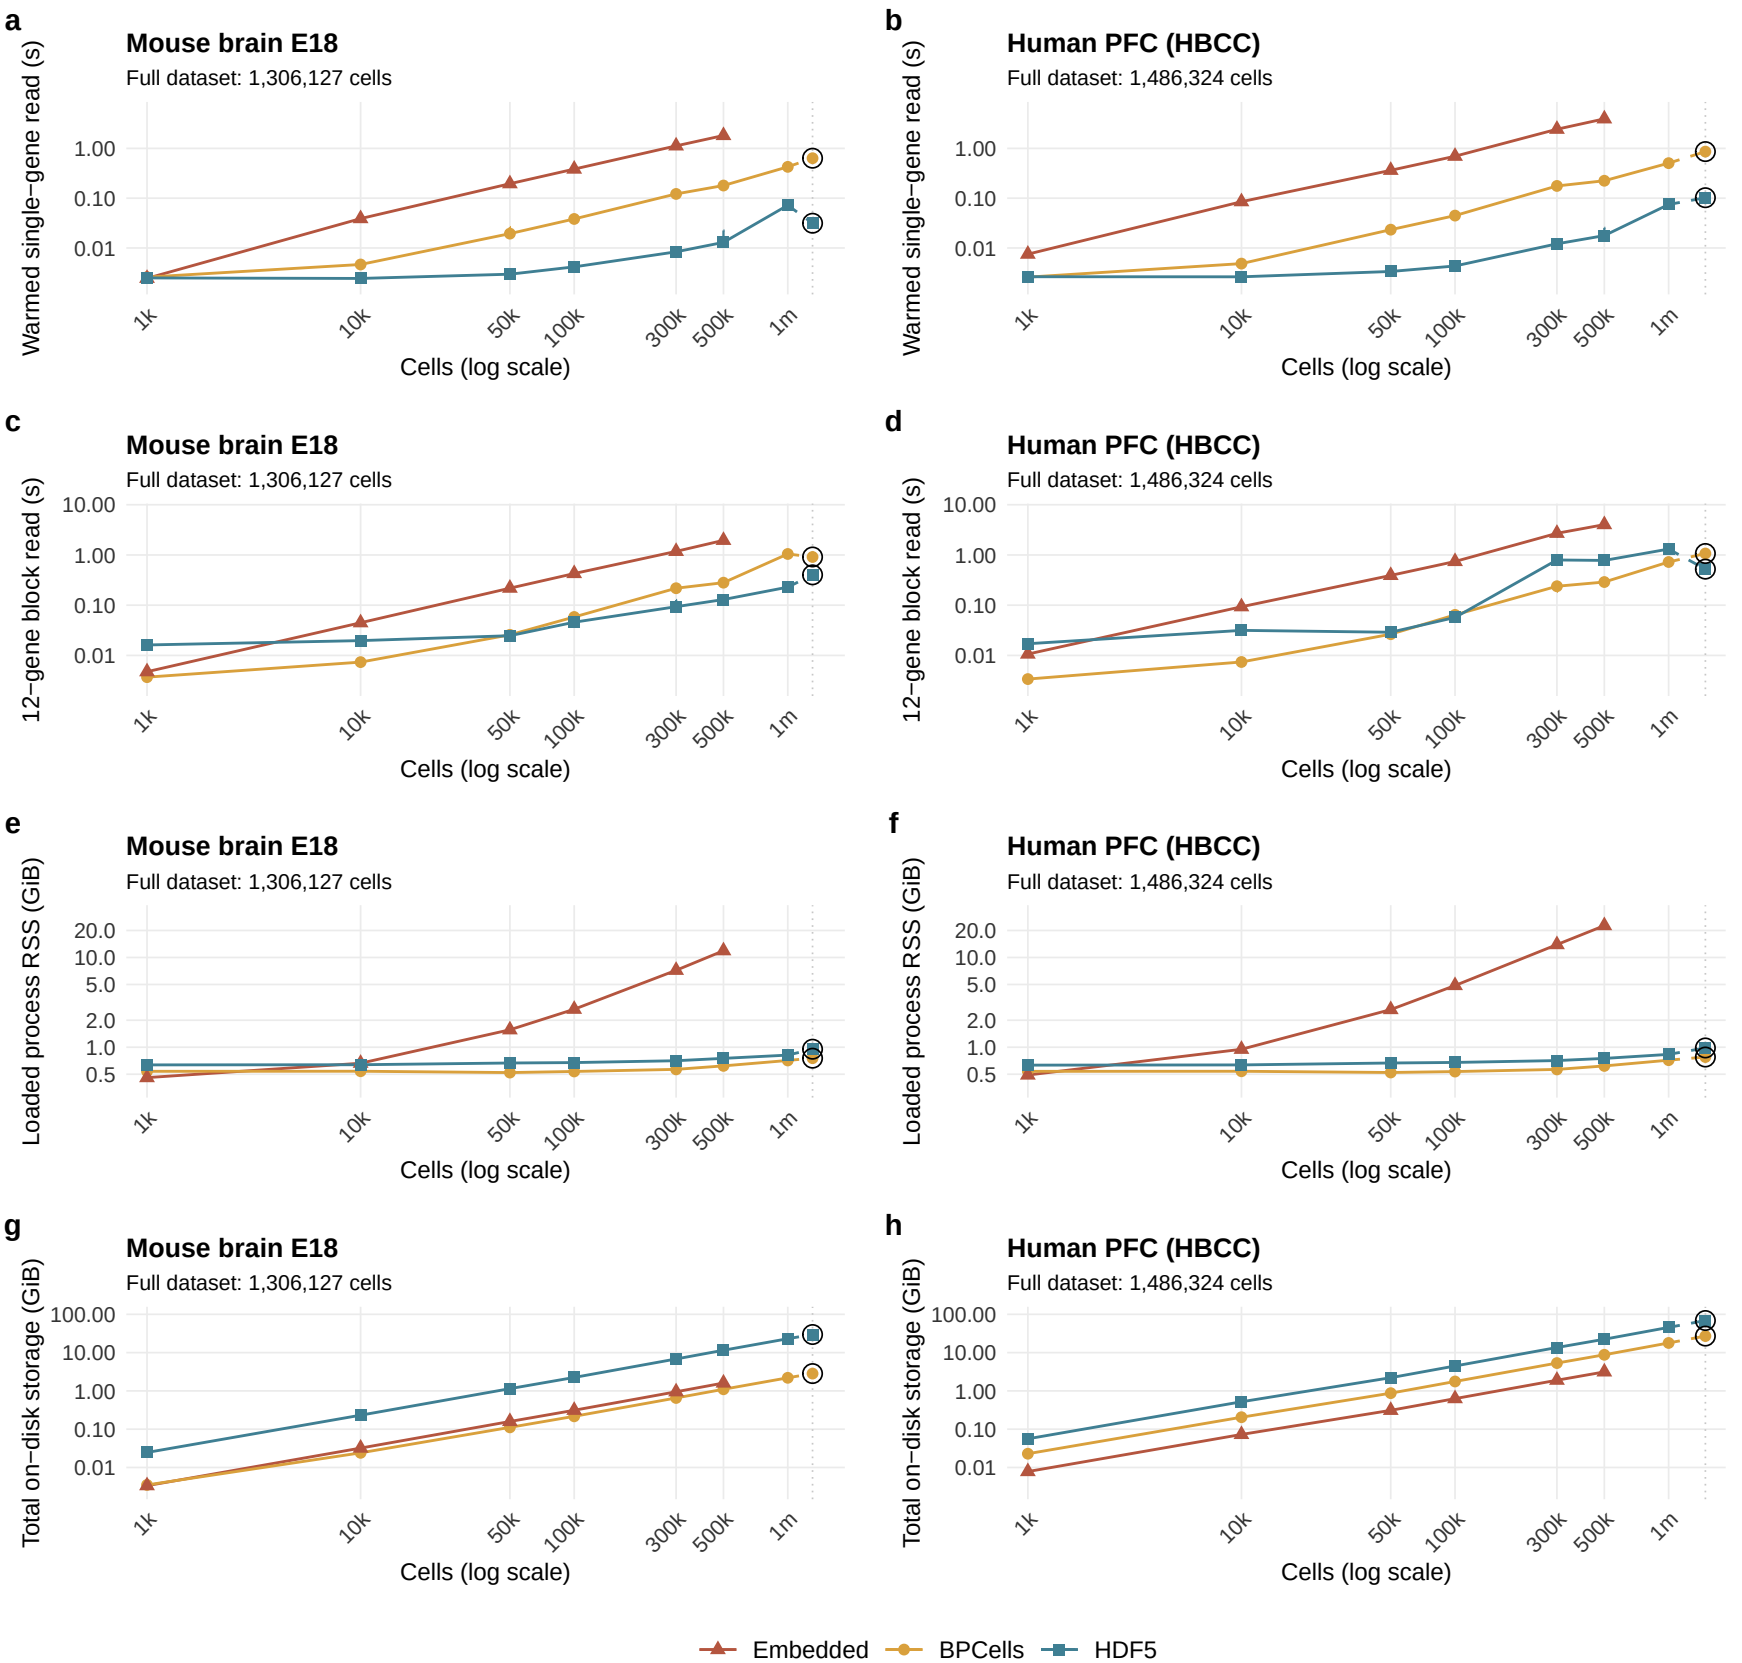

Both axes are logarithmic. Rings/dashed extensions: reused Full dataset measurements. 300k and Human 500k repeats: new measurements.
Filled points: median across 5 builds; whiskers: min−max. Access: median of 2 processes within each build.
Human embedded 500k: five builds (one feasibility plus four prospective repeats) under 99%−of−available caps.
Same query strategy at every tier; selected genes vary with tier. Lines join observations, not fitted scaling laws.
Original Human embedded 500k resource skips retained. Embedded 1m/full not scheduled. RSS is after load.