Expression backend scaling through the full datasets

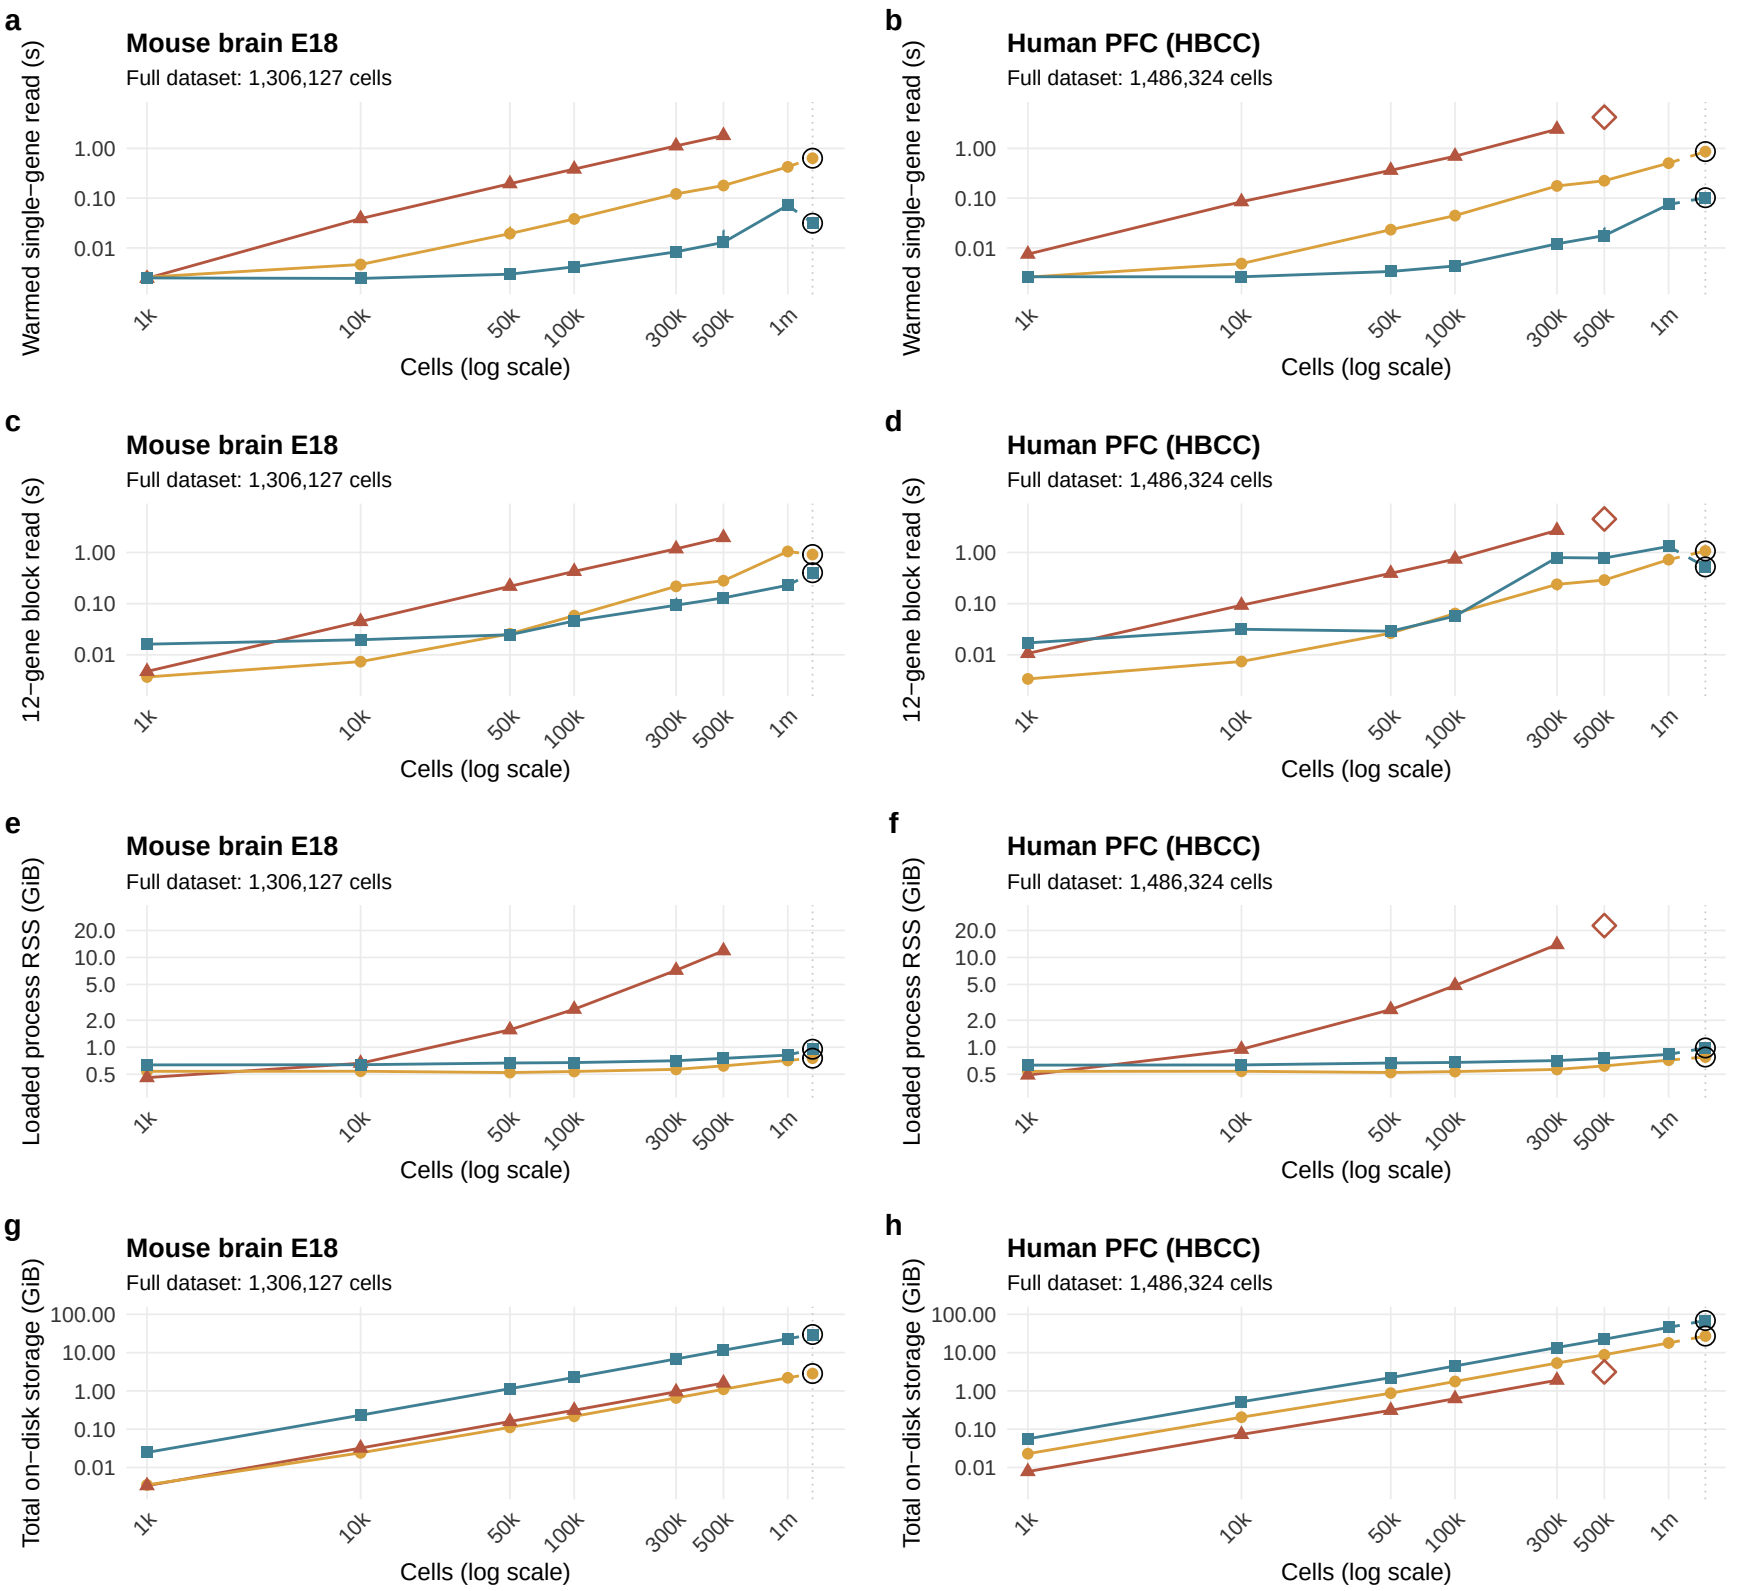

Both axes are logarithmic. Rings/dashed extensions: reused Full dataset measurements. 300k: new measurements.
Filled points: median across 5 builds; whiskers: min−max. Access: median of 2 processes within each build.
Hollow diamonds: Human embedded 500k, 1 build and 2 reads under a 99 GiB cap; no across−build uncertainty.
Same query strategy at every tier; selected genes vary with tier. Lines join observations, not fitted scaling laws.
Original Human embedded 500k resource skips retained. Embedded 1m/full not scheduled. RSS is after load.